

Models

CS109a: Fall 2018

Authors: Gordon Hew, Wenqin Hu, Blair Leduc

TF: Ken Arnold

```
In [1]: %matplotlib inline
import numpy as np
import numpy.random as nd
import pandas as pd
import math
import matplotlib.pyplot as plt
import json
import matplotlib
import gzip
import os
import seaborn as sns

from sklearn.linear_model import LogisticRegressionCV
from sklearn.linear_model import LogisticRegression
from sklearn.model_selection import cross_val_score
from sklearn.metrics import accuracy_score
from sklearn.metrics import confusion_matrix
from sklearn.preprocessing import Imputer
from sklearn.metrics import mean_squared_error
from sklearn.model_selection import train_test_split
from IPython.display import display
from pandas.plotting import scatter_matrix

from sklearn.impute import SimpleImputer
from sklearn.model_selection import train_test_split
from sklearn import preprocessing
from sklearn.ensemble import RandomForestClassifier
from sklearn.utils import resample
from sklearn.tree import DecisionTreeClassifier

from scipy.stats import ks_2samp

import statsmodels.api as sm
from statsmodels.api import OLS
```

Set botometer score cutoff for non-bot and bot

load botometer scores from file

```

In [2]: def arrange_botometer_score(data_botometer, is_bot):
    account_score = []
    for i in (range(len(data_botometer))):
        for key, val in data_botometer[i].items():
            temp = key
            if temp != 'error':
                parameters = []
                parameters.append(is_bot)
                for key, val in data_botometer[i]['user'].items():
                    parameters.append(val)
                for key, val in data_botometer[i]['categories'].items():
                    parameters.append(val)
                for key, val in data_botometer[i]['scores'].items():
                    parameters.append(val)
                account_score.append(parameters)

    account_score = pd.DataFrame(account_score)
    account_score.columns = columns = ['is_bot', 'id', 'screen_name',
                                        'content', 'friend', 'network',
                                        'sentiment', 'temporal', 'user',
                                        'score_english', 'score_universal'
                                        ]

    return account_score

```

```

In [3]: with open('data/botometer_result_1000random.json') as handle:
    random_botometer = json.loads(handle.read())
    with open('data/botometer_result_936bot.json') as handle:
        bot_botometer = json.loads(handle.read())
    with open('data/botometer_result_1000_verified.json') as handle:
        human_botometer = json.loads(handle.read())

```

```

In [4]: human_score = arrange_botometer_score(human_botometer, is_bot='False')
    bot_score = arrange_botometer_score(bot_botometer, is_bot='Ture')
    random_score = arrange_botometer_score(random_botometer, is_bot='unknow')

```

```

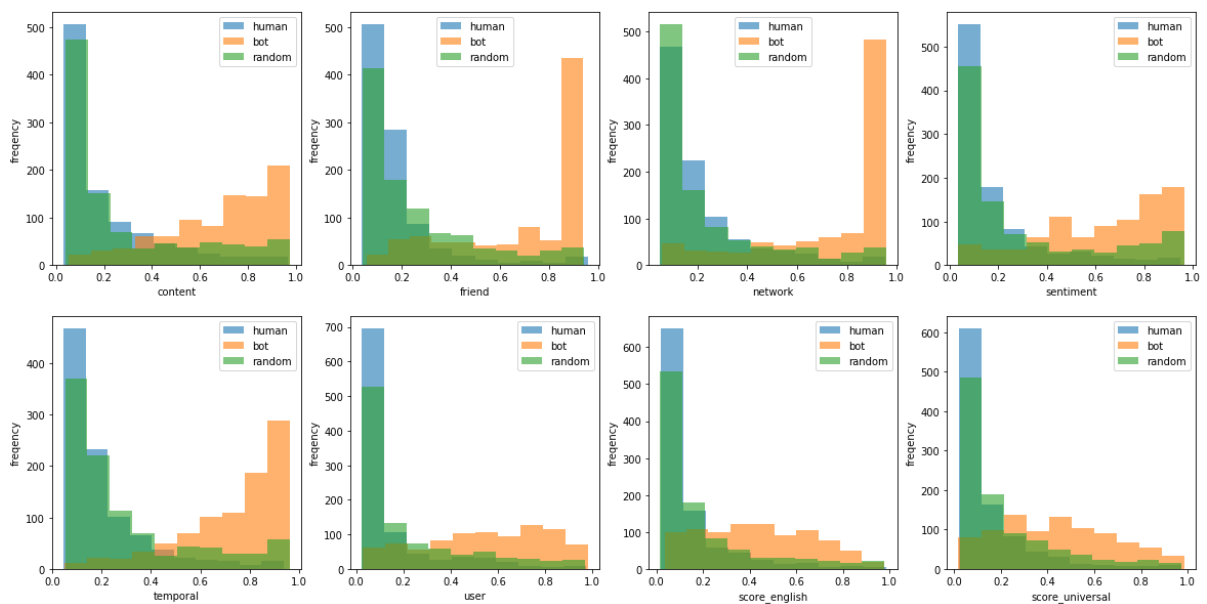
In [5]: fig, axs = plt.subplots(2, 4, figsize=(20, 10))
col = 3
for j in range(2):
    for i in range(4):
        axs[j,i].hist(human_score.iloc[:,col].tolist(), alpha=0.6,
                        label='human')
        axs[j,i].hist(bot_score.iloc[:,col].tolist(), alpha=0.6,
                        label='bot')
        axs[j,i].hist(random_score.iloc[:,col].tolist(), alpha=0.6,
                        label='random')
        axs[j,i].set_ylabel('frequency')
        axs[j,i].set_xlabel(human_score.columns[col])
        axs[j,i].legend()
        col=col+1

fig.suptitle("Botometer scores on 936 botwiki, 1000 human "\
            +"and 1000 random twitter users", fontsize=16)

fig.savefig('fig/Botometer scores distribution.png', format = 'png')

```

Botometer scores on 936 botwiki, 1000 human and 1000 random twitter users



Find the ratio of non-bot and bot in random set


```

In [6]: bins = np.arange(0,1.2,0.1)
plt.figure(figsize=(15, 8))
plt.hist([human_score.iloc[:,10].tolist(),
         bot_score.iloc[:,10].tolist(),
         random_score.iloc[:,10].tolist()],
         alpha=0.5, bins=bins, color=['C0', 'C3', 'C2'],
         label = ['Non Bot', 'Bot', 'Random'])
#plt.hist(, alpha=0.5, bins=bins, color='C3', label = 'Bot')
#plt.hist(, alpha=0.5, bins=bins, color='C2', label = 'Random')
plt.ylabel('Frequency', fontsize = 18)
plt.xlabel('Botometer Universal Score')
plt.xticks(np.arange(-0.05, 1.1, 0.1))
plt.xlim([0,1.1])
plt.grid(axis='y', linestyle='--')
plt.legend()

plt.suptitle("Botometer Universal Score Distribution", fontsize = 18)
plt.savefig('fig/score_universal_distribution.png', format = 'png')
plt.savefig('../docs/assets/images/score_universal_distribution.png',
            format = 'png')

```

Botometer Universal Score Distribution



```
In [7]: np.random.seed(123456)
x = random_score['score_universal']
y = human_score['score_universal']
z = bot_score['score_universal']

print('difference between random and human')
print(ks_2samp(x, y))
print('difference between random and bot')
print(ks_2samp(x, z))
```

difference between random and human
Ks_2sampResult(statistic=0.1559714867617108, pvalue=5.175986131548254e-11)
difference between random and bot
Ks_2sampResult(statistic=0.47200000000000003, pvalue=1.3333627704181822e-92)

```
In [8]: np.random.seed(123)

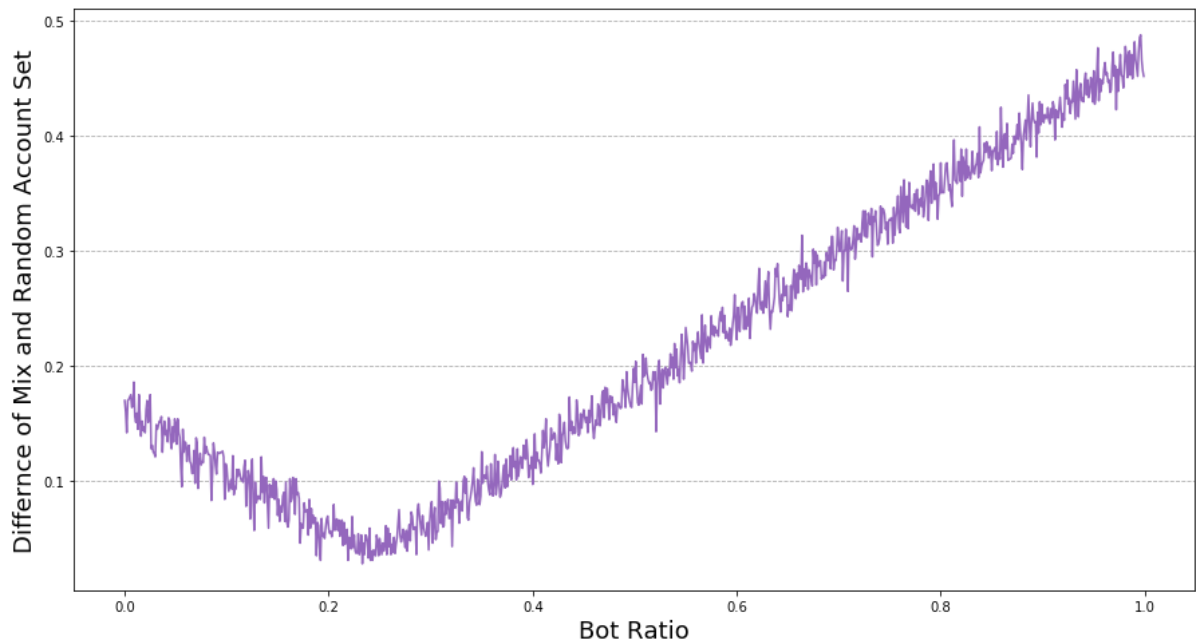
def loss_function_bot_ratio(bot_ratio):
    human_number = int(len(x) * (1-bot_ratio))
    bot_number = int(len(x) * bot_ratio)
    mix = np.random.choice(y, human_number).tolist() \
          + np.random.choice(z, bot_number).tolist()
    return ks_2samp(x, mix)[0]
```

```
In [9]: p_random_mix = []
for bot_ratio in np.arange(0, 1, 0.001):
    p = loss_function_bot_ratio(bot_ratio)
    p_random_mix.append(p)
```

```
In [10]: plt.figure(figsize=(15, 8))
plt.plot(np.arange(0, 1, 0.001), p_random_mix, color='C4')
plt.xlabel("Bot Ratio", fontsize = 18)
plt.ylabel("Difference of Mix and Random Account Set", fontsize = 18)
plt.grid(axis='y', linestyle='--')

plt.suptitle("Bot Ratio vs Difference of Mix and Random Account Set",
            fontsize = 18)
plt.savefig('fig/bot_ratio.png', format = 'png')
plt.savefig('../docs/assets/images/bot_ratio.png', format = 'png')
```

Bot Ratio vs Difference of Mix and Random Account Set



```
In [11]: best_bot_ratio = np.arange(0, 1, 0.001)\
            [p_random_mix.index(min(p_random_mix))]
print("ratio of bot in random account set: ", best_bot_ratio)

ratio of bot in random account set: 0.233
```

Determin botmeter score cutoff : logistic regression

```
In [12]: len(human_score.sample(int (len (x) * best_bot_ratio)))
```

Out[12]: 233

```
In [13]: best_human_number = int(len(x)*(1-best_bot_ratio))
best_bot_number = int(len(x)*best_bot_ratio)

human_bot_score_df = human_score.sample(best_human_number)\
                    .append(bot_score.sample(best_bot_number))
len (human_bot_score_df)
print("training data including mixing of:")
print("verified human: ",best_human_number,
      " bot:",best_bot_number,
      " total: ",len(human_bot_score_df))
```

```
training data including mixing of:
verified human: 767 bot: 233 total: 1000
```

```
In [14]: X_data = human_bot_score_df.iloc[:,10:11]
y_data = pd.get_dummies(human_bot_score_df['is_bot'],
                        drop_first=True)['Ture']

X_train, X_test = train_test_split(X_data, test_size = 0.2,
                                   random_state=90)
y_train, y_test = train_test_split(y_data, test_size = 0.2,
                                   random_state=90)
```

```
In [15]: def perfromance(model,X_train,X_test):
y_pred_train = model.predict(X_train)
y_pred_test = model.predict(X_test)

#Perfromance Evaluation
train_score = accuracy_score(y_train, y_pred_train)*100
test_score = accuracy_score(y_test, y_pred_test)*100

print("logistic regression model")
print("Training Set Accuracy:",str(train_score)+'%')
print("Testing Set Accuracy:",str(test_score)+'%')

print()
print("Rows: True Lables (0,1), \nColumnms: Predicted Lables (0,1)")
print("human", "bot")

print("train:")
print(confusion_matrix(y_train, y_pred_train))
print("test:")
print(confusion_matrix(y_test, y_pred_test))
```

```
In [16]: botometer_train_model = LogisticRegressionCV(Cs=10, cv=5)\
                                                .fit(X_train, y_train)
perfomance(botometer_train_model,X_train,X_test)
```

```
logistic regression model
Training Set Accuracy: 83.375%
Testing Set Accuracy: 79.0%

Rows: True Lables (0,1),
Columns: Predicted Lables (0,1)
human bot
train:
[[582  35]
 [ 98  85]]
test:
[[141   9]
 [ 33  17]]
```

[illegible]

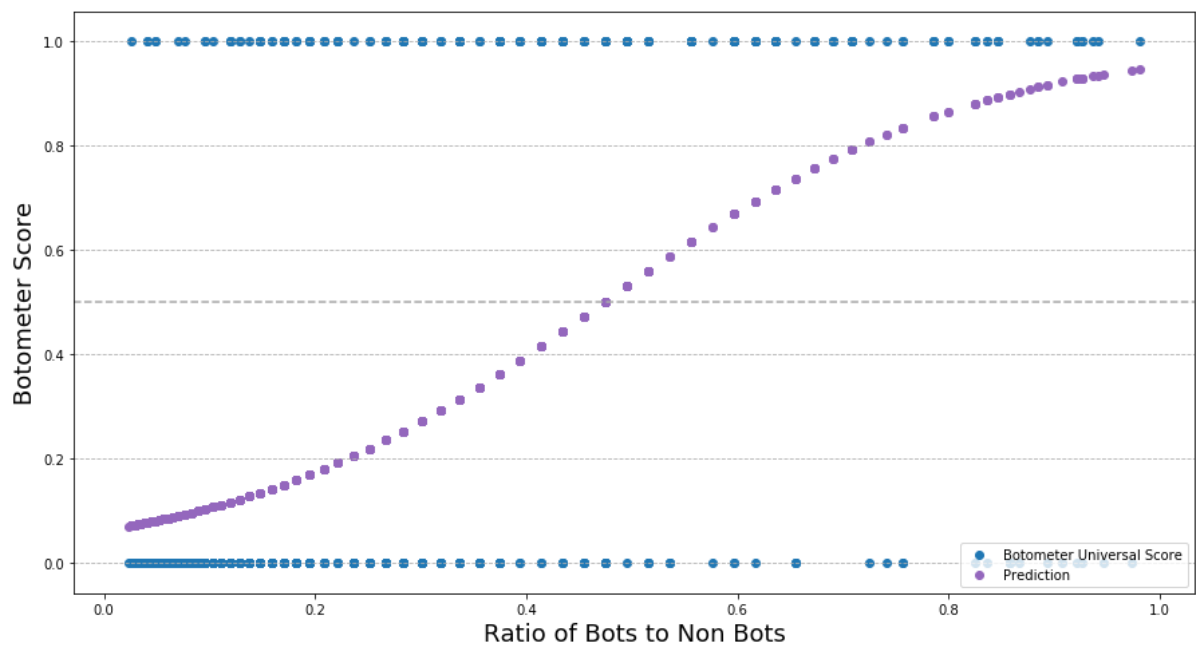

```

In [18]: plt.figure(figsize=(15, 8))
plt.scatter(X_train, y_train, label = "Botometer Universal Score",
            color='C0')
plt.scatter(X_train, botometer_train_model.predict_proba(X_train)[: ,1],
            label = "Prediction", color='C4')
plt.axhline(y=0.5, color='darkgray', linestyle='--')
plt.xlabel("Ratio of Bots to Non Bots", fontsize = 18)
plt.ylabel("Botometer Score", fontsize = 18)
plt.legend(loc='lower right')
plt.grid(axis='y', linestyle='--')

plt.suptitle("Botometer Score vs Ratio of Bots to Non Bots",
            fontsize = 18)
plt.savefig('fig/bot_cutoff.png', format = 'png')
plt.savefig('../docs/assets/images/bot_cutoff.png', format = 'png')

```

Botometer Score vs Ratio of Bots to Non Bots



```

In [19]: diff = (abs(botometer_train_model.predict_proba(X_train)[: ,1]-0.5))\
              .tolist()
cut_off = X_train["score_universal"].tolist()[diff.index(min(diff))]

```

```

In [20]: print ("best cut off of score_universal_botometer is", round(cut_off,2))

best cut off of score_universal_botometer is 0.47

```

Read and clear data (1000 random accounts)


```
In [22]: pd.set_option('display.max_rows', 350)
pd.set_option('display.max_colwidth', -1)

print(users_summary_df.shape)
display(users_summary_df.describe())
display(users_summary_df.dtypes)
```

(1000, 95)

| | favourites_count | followers_count | friends_count | listed_count | statuses_count | scores.unive |
|--------------|------------------|-----------------|---------------|--------------|----------------|--------------|
| count | 1000.000000 | 1000.000000 | 1000.000000 | 1000.000000 | 1000.000000 | 997.000 |
| mean | 21331.004000 | 1867.186000 | 1139.991000 | 25.115000 | 31638.967000 | 0.214 |
| std | 39114.780324 | 7344.047951 | 4073.343137 | 92.364381 | 59896.719907 | 0.214 |
| min | 0.000000 | 0.000000 | 0.000000 | 0.000000 | 1.000000 | 0.024 |
| 25% | 1550.250000 | 108.750000 | 167.000000 | 0.000000 | 2309.750000 | 0.064 |
| 50% | 7982.500000 | 363.000000 | 397.500000 | 2.000000 | 9648.500000 | 0.114 |
| 75% | 23129.750000 | 1094.500000 | 908.250000 | 11.000000 | 32057.500000 | 0.304 |
| max | 383288.000000 | 100730.000000 | 85123.000000 | 1550.000000 | 624250.000000 | 0.974 |

8 rows × 56 columns

| | |
|------------------------------------|---------|
| contributors_enabled | bool |
| created_at | object |
| default_profile | bool |
| default_profile_image | bool |
| description | object |
| entities.description.urls | object |
| entities.url.urls | object |
| favourites_count | int64 |
| follow_request_sent | bool |
| followers_count | int64 |
| following | bool |
| friends_count | int64 |
| geo_enabled | bool |
| has_extended_profile | bool |
| id_str | object |
| is_translation_enabled | bool |
| is_translator | bool |
| lang | object |
| listed_count | int64 |
| location | object |
| name | object |
| notifications | bool |
| profile_background_color | object |
| profile_background_image_url | object |
| profile_background_image_url_https | object |
| profile_background_tile | bool |
| profile_banner_url | object |
| profile_image_url | object |
| profile_image_url_https | object |
| profile_link_color | object |
| profile_sidebar_border_color | object |
| profile_sidebar_fill_color | object |
| profile_text_color | object |
| profile_use_background_image | bool |
| protected | bool |
| screen_name | object |
| statuses_count | int64 |
| time_zone | object |
| translator_type | object |
| url | object |
| utc_offset | object |
| verified | bool |
| user.id_str | object |
| scores.universal | float64 |
| tweets_per_hour_00 | float64 |
| tweets_per_hour_01 | float64 |
| tweets_per_hour_02 | float64 |
| tweets_per_hour_03 | float64 |
| tweets_per_hour_04 | float64 |
| tweets_per_hour_05 | float64 |
| tweets_per_hour_06 | float64 |
| tweets_per_hour_07 | float64 |
| tweets_per_hour_08 | float64 |
| tweets_per_hour_09 | float64 |
| tweets_per_hour_10 | float64 |
| tweets_per_hour_11 | float64 |
| tweets_per_hour_12 | float64 |

| | |
|-------------------------------------|---------|
| tweets_per_hour_13 | float64 |
| tweets_per_hour_14 | float64 |
| tweets_per_hour_15 | float64 |
| tweets_per_hour_16 | float64 |
| tweets_per_hour_17 | float64 |
| tweets_per_hour_18 | float64 |
| tweets_per_hour_19 | float64 |
| tweets_per_hour_20 | float64 |
| tweets_per_hour_21 | float64 |
| tweets_per_hour_22 | float64 |
| tweets_per_hour_23 | float64 |
| tweets_per_hour | object |
| mean_links_per_tweet | float64 |
| mean_words_per_tweet | float64 |
| mean_hashtags_per_tweet | float64 |
| mean_user_mentions_per_tweet | float64 |
| mean_favourites_per_tweet | float64 |
| mean_media_per_tweet | float64 |
| mean_user_symbols_per_tweet | float64 |
| mean_retweets_per_tweet | float64 |
| mean_truncations_per_tweet | float64 |
| mean_links_to_twitter | float64 |
| mean_links_to_top_social_media | float64 |
| mean_links_to_top_digital_media | float64 |
| mean_links_to_top_news | float64 |
| mean_links_to_top_products_services | float64 |
| mean_links_to_top_celebrities | float64 |
| mean_links_to_top_organizations | float64 |
| mean_links_to_top_sports | float64 |
| mean_links_to_top_adult | float64 |
| retweet_ratio | float64 |
| mean_ref_to_person | float64 |
| mean_ref_to_norp | float64 |
| mean_ref_to_org | float64 |
| mean_ref_to_gpe | float64 |
| mean_ref_to_product | float64 |
| mean_ref_to_law | float64 |
| mean_ref_to_money | float64 |
| dtype: | object |

```
In [23]: users = users_summary_df.copy()
```

```
In [24]: type_bool = []
type_int = []
type_float = []
type_other = []
for i in range(len(users.dtypes)):
    if users.dtypes[i] == 'bool':
        type_bool.append(i)
    if users.dtypes[i] == 'int64':
        type_int.append(i)
    if users.dtypes[i] == 'float64':
        type_float.append(i)
    if users.dtypes[i] != 'bool' \
        and users.dtypes[i] != 'int64' \
        and users.dtypes[i] != 'float64':
        type_other.append(i)
print (len(type_bool),len(type_int),len(type_float),len(type_other))
```

14 5 51 25

```

In [25]: users_clear = users.iloc[:,([43,14,20,7, 9, 11, 18, 36] \
                                     + type_float[1:]+[2,3,12,13,15,25,33,41])]
users_clear.index = range(1000)
print(users_clear.shape)
print(users_clear.columns)

(1000, 66)
Index(['scores.universal', 'id_str', 'name', 'favourites_count',
      'followers_count', 'friends_count', 'listed_count', 'statuses_co
unt',
      'tweets_per_hour_00', 'tweets_per_hour_01', 'tweets_per_hour_0
2',
      'tweets_per_hour_03', 'tweets_per_hour_04', 'tweets_per_hour_0
5',
      'tweets_per_hour_06', 'tweets_per_hour_07', 'tweets_per_hour_0
8',
      'tweets_per_hour_09', 'tweets_per_hour_10', 'tweets_per_hour_1
1',
      'tweets_per_hour_12', 'tweets_per_hour_13', 'tweets_per_hour_1
4',
      'tweets_per_hour_15', 'tweets_per_hour_16', 'tweets_per_hour_1
7',
      'tweets_per_hour_18', 'tweets_per_hour_19', 'tweets_per_hour_2
0',
      'tweets_per_hour_21', 'tweets_per_hour_22', 'tweets_per_hour_2
3',
      'mean_links_per_tweet', 'mean_words_per_tweet',
      'mean_hashtags_per_tweet', 'mean_user_mentions_per_tweet',
      'mean_favourites_per_tweet', 'mean_media_per_tweet',
      'mean_user_symbols_per_tweet', 'mean_retweets_per_tweet',
      'mean_truncations_per_tweet', 'mean_links_to_twitter',
      'mean_links_to_top_social_media', 'mean_links_to_top_digital_med
ia',
      'mean_links_to_top_news', 'mean_links_to_top_products_services',
      'mean_links_to_top_celebrities', 'mean_links_to_top_organization
s',
      'mean_links_to_top_sports', 'mean_links_to_top_adult', 'retweet_
ratio',
      'mean_ref_to_person', 'mean_ref_to_norp', 'mean_ref_to_org',
      'mean_ref_to_gpe', 'mean_ref_to_product', 'mean_ref_to_law',
      'mean_ref_to_money', 'default_profile', 'default_profile_image',
      'geo_enabled', 'has_extended_profile', 'is_translation_enabled',
      'profile_background_tile', 'profile_use_background_image', 'veri
fied'],
      dtype='object')

```



```
In [26]: X_data = users_clear.iloc[:,3:]
y_data = []
bot_number = 0
for i in range(len(users_clear)):
    if users_clear['scores.universal'][i] >= BOTOMETER_SCORE_THRESHOLD:
        y_data = y_data + [1]
        bot_number = bot_number + 1
    else: y_data = y_data + [0]
print(X_data.shape, len(y_data))
print("bot number:" ,bot_number)

(1000, 63) 1000
bot number: 149
```

```
In [27]: X_train, X_test = train_test_split(X_data,
                                             test_size = 0.2,
                                             random_state=90)
y_train, y_test = train_test_split(y_data,
                                    test_size = 0.2,
                                    random_state=90)
```

Logistic Regresion

```

In [28]: def perfromance(model,X_train,X_test):
    y_pred_train = model.predict(X_train)
    y_pred_test = model.predict(X_test)

    #Perfromance Evaluation
    train_score = accuracy_score(y_train, y_pred_train)*100
    train_matrix = confusion_matrix(y_train, y_pred_train)
    test_score = accuracy_score(y_test, y_pred_test)*100
    test_matrix = confusion_matrix(y_test, y_pred_test)

    train_ture_negative = train_matrix[0][0]\
                        /(train_matrix[0][0]+train_matrix[0][1])*100
    train_ture_positive = train_matrix[1][1]\
                        /(train_matrix[1][0]+train_matrix[1][1])*100

    test_ture_negative = test_matrix[0][0]\
                        /(test_matrix[0][0]+test_matrix[0][1])*100
    test_ture_positive = test_matrix[1][1]\
                        /(test_matrix[1][0]+test_matrix[1][1])*100

    print()
    print("Rows: True Lables (0,1), \nColumns: Predicted Lables (0,1)")
    print("human", "bot")

    print("train:")
    print(train_matrix)
    print("test:")
    print(test_matrix)

    print("Training Set total accuracy:",str(train_score)+'%')
    print("Training Set non-bot accuracy:",str(train_ture_negative)+'%')

    print("Training Set bot accuracy:",str(train_ture_positive)+'%')

    print("Test Set accuracy:",str(test_score)+'%')
    print("Test Set non-bot accuracy:",str(test_ture_negative)+'%')
    print("Test Set bot accuracy:",str(test_ture_positive)+'%')

```

```

In [29]: def tuning_bot_weight(bot_weight):
            model = LogisticRegressionCV(Cs=10, cv=5,max_iter=1000,
                                         class_weight = {0:1,1:bot_weight})\
                                         .fit(X_train,
                                              y_train)

            y_pred_train = model.predict(X_train)
            y_pred_test = model.predict(X_test)

            train_score = accuracy_score(y_train, y_pred_train)*100
            train_matrix = confusion_matrix(y_train, y_pred_train)
            test_score = accuracy_score(y_test, y_pred_test)*100
            test_matrix = confusion_matrix(y_test, y_pred_test)

            train_ture_negative = train_matrix[0][0]\
                                /(train_matrix[0][0]+train_matrix[0][1])*100
            train_ture_positive = train_matrix[1][1]\
                                /(train_matrix[1][0]+train_matrix[1][1])*100

            test_ture_negative = test_matrix[0][0]\
                                /(test_matrix[0][0]+test_matrix[0][1])*100
            test_ture_positive = test_matrix[1][1]\
                                /(test_matrix[1][0]+test_matrix[1][1])*100

            accuracy = [bot_weight, train_score, train_ture_negative,
                        train_ture_positive, test_score, test_ture_negative,
                        test_ture_positive]
            return accuracy

```

```

In [30]: bot_weight_accuracy = []
            for bot_weight in range(5,11):
                bot_weight_accuracy.append(tuning_bot_weight(bot_weight))
            bot_weight_accuracy = pd.DataFrame(bot_weight_accuracy,
                                              columns =['bot_weight',
                                                      'train_score',
                                                      'train_ture_negative',
                                                      'train_ture_positive',
                                                      'test_score',
                                                      'test_ture_negative',
                                                      'test_ture_positive'] )

```

/Users/blair/.pyenv/versions/3.6.7/lib/python3.6/site-packages/sklearn/linear_model/logistic.py:757: ConvergenceWarning: lbfgs failed to converge. Increase the number of iterations.
 "of iterations.", ConvergenceWarning)

```

In [31]: fig, axs = plt.subplots(1, 2, figsize=(15, 5))

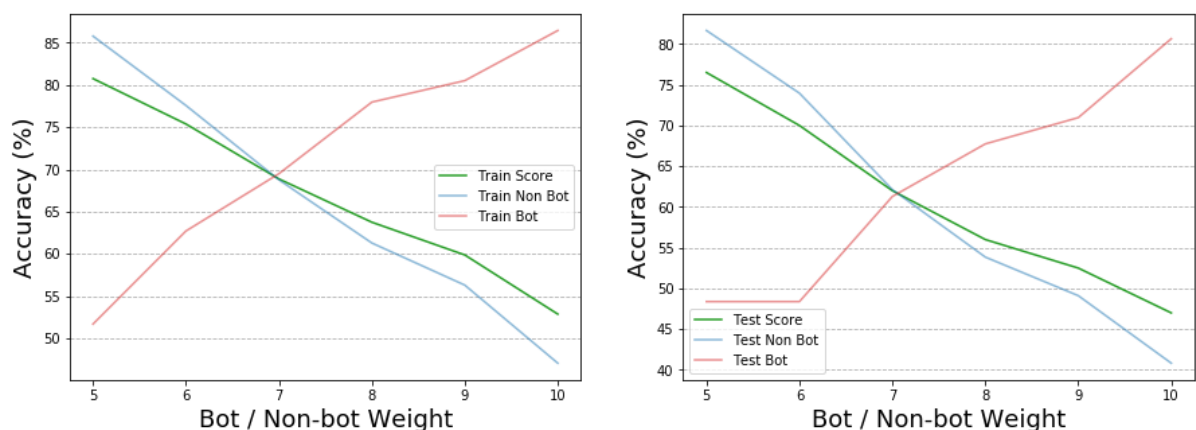
axs[0].plot(bot_weight_accuracy.iloc[:,0],bot_weight_accuracy.iloc[:,1],
            label='Train Score', color='C2')
axs[0].plot(bot_weight_accuracy.iloc[:,0],bot_weight_accuracy.iloc[:,2],
            label='Train Non Bot', color='C0', alpha=0.5)
axs[0].plot(bot_weight_accuracy.iloc[:,0],bot_weight_accuracy.iloc[:,3],
            label='Train Bot', color='C3', alpha=0.5)
axs[0].set_xlabel('Bot / Non-bot Weight', fontsize = 18)
axs[0].set_ylabel('Accuracy (%)', fontsize = 18)
axs[0].legend()
axs[0].grid(axis='y', linestyle='--')

axs[1].plot(bot_weight_accuracy.iloc[:,0],bot_weight_accuracy.iloc[:,4],
            label='Test Score', color='C2')
axs[1].plot(bot_weight_accuracy.iloc[:,0],bot_weight_accuracy.iloc[:,5],
            label='Test Non Bot', color='C0', alpha=0.5)
axs[1].plot(bot_weight_accuracy.iloc[:,0],bot_weight_accuracy.iloc[:,6],
            label='Test Bot', color='C3', alpha=0.5)
axs[1].set_xlabel('Bot / Non-bot Weight', fontsize = 18)
axs[1].set_ylabel('Accuracy (%)', fontsize = 18)
axs[1].legend()
axs[1].grid(axis='y', linestyle='--')

fig.suptitle("Logistic Regression Model - Accuracy vs Weights",
            fontsize=18)
plt.savefig('../docs/assets/images/log_regression_model.png',
            format = 'png')
fig.savefig('fig/log_regression_model.png', format = 'png')

```

Logistic Regression Model - Accuracy vs Weights



```
In [32]: LR_model = LogisticRegressionCV(Cs=10, cv=5,max_iter=1000,
                                             class_weight = {0:1,1:7}).fit(X_train,
                                                                           y_train)

perfomance(LR_model,X_train,X_test)

Rows: True Lables (0,1),
Columnns: Predicted Lables (0,1)
human bot
train:
[[469 213]
 [ 36  82]]
test:
[[105  64]
 [ 12  19]]
Training Set total accuracy: 68.875%
Training Set non-bot accuracy: 68.76832844574781%
Training Set bot accuracy: 69.49152542372882%
Test Set accuracy: 62.0%
Test Set non-bot accuracy: 62.1301775147929%
Test Set bot accuracy: 61.29032258064516%
```

Random Forest

```
In [33]: n_trees = 50
tree_depth = 5
RF_model = RandomForestClassifier(n_estimators=n_trees,
                                  max_depth=tree_depth,
                                  class_weight = {0:1,1:20}).fit(X_train,
                                                                y_train)
```

```
In [34]: perfomance(RF_model,X_train,X_test)

Rows: True Lables (0,1),
Columnns: Predicted Lables (0,1)
human bot
train:
[[589  93]
 [  2 116]]
test:
[[136  33]
 [  4  27]]
Training Set total accuracy: 88.125%
Training Set non-bot accuracy: 86.36363636363636%
Training Set bot accuracy: 98.30508474576271%
Test Set accuracy: 81.5%
Test Set non-bot accuracy: 80.4733727810651%
Test Set bot accuracy: 87.09677419354838%
```

```
In [35]: def tuning_tree_depth(tree_depth):
    model = RandomForestClassifier(n_estimators=n_trees,
                                  max_depth=tree_depth,
                                  class_weight = {0:1,1:20}).fit(X_train,
                                                                  y_train)

    y_pred_train = model.predict(X_train)
    y_pred_test = model.predict(X_test)

    train_score = accuracy_score(y_train, y_pred_train)*100
    train_matrix = confusion_matrix(y_train, y_pred_train)
    test_score = accuracy_score(y_test, y_pred_test)*100
    test_matrix = confusion_matrix(y_test, y_pred_test)

    train_ture_negative = train_matrix[0][0]\
                          /(train_matrix[0][0]+train_matrix[0][1])*100
    train_ture_positive = train_matrix[1][1]\
                          /(train_matrix[1][0]+train_matrix[1][1])*100

    test_ture_negative = test_matrix[0][0]\
                        /(test_matrix[0][0]+test_matrix[0][1])*100
    test_ture_positive = test_matrix[1][1]\
                        /(test_matrix[1][0]+test_matrix[1][1])*100

    accuracy = [tree_depth, train_score, train_ture_negative,
                train_ture_positive, test_score, test_ture_negative,
                test_ture_positive]
    return accuracy
```

```
In [36]: tree_depth_accuracy = []
    for tree_depth in range(1,30):
        tree_depth_accuracy.append(tuning_tree_depth(tree_depth))
    tree_depth_accuracy = pd.DataFrame(tree_depth_accuracy,
                                       columns =['tree_depth',
                                                'train_score',
                                                'train_ture_negative',
                                                'train_ture_positive',
                                                'test_score',
                                                'test_ture_negative',
                                                'test_ture_positive'] )
```

```
In [37]: tree_depth_accuracy.head()
```

Out[37]:

| | tree_depth | train_score | train_ture_negative | train_ture_positive | test_score | test_ture_negative |
|---|------------|-------------|---------------------|---------------------|------------|--------------------|
| 0 | 1 | 14.750 | 0.000000 | 100.0 | 15.5 | 0.000000 |
| 1 | 2 | 18.250 | 4.105572 | 100.0 | 17.0 | 1.775148 |
| 2 | 3 | 48.250 | 39.296188 | 100.0 | 48.5 | 39.053254 |
| 3 | 4 | 72.375 | 67.595308 | 100.0 | 67.5 | 62.130178 |
| 4 | 5 | 86.750 | 84.457478 | 100.0 | 79.5 | 78.106509 |

```

In [38]: fig, axs = plt.subplots(1, 2, figsize=(15, 5))

axs[0].plot(tree_depth_accuracy.iloc[:,0],tree_depth_accuracy.iloc[:,1],

            label='Train Score', color='C2')
axs[0].plot(tree_depth_accuracy.iloc[:,0],tree_depth_accuracy.iloc[:,2],

            label='Train Non Bot', color='C0', alpha=0.5)
axs[0].plot(tree_depth_accuracy.iloc[:,0],tree_depth_accuracy.iloc[:,3],

            label='Train Bot', color='C3', alpha=0.5)
axs[0].set_xlabel('Tree Depth', fontsize=18)
axs[0].set_ylabel('Accuracy (%)', fontsize=18)
axs[0].legend()
axs[0].grid(axis='y', linestyle='--')

axs[1].plot(tree_depth_accuracy.iloc[:,0],tree_depth_accuracy.iloc[:,4],


            label='Test Score', color='C2')
axs[1].plot(tree_depth_accuracy.iloc[:,0],tree_depth_accuracy.iloc[:,5],

            label='Test Non Bot', color='C0', alpha=0.5)
axs[1].plot(tree_depth_accuracy.iloc[:,0],tree_depth_accuracy.iloc[:,6],

            label='Test Bot', color='C3', alpha=0.5)
axs[1].set_xlabel('Tree Depth', fontsize=18)
axs[1].set_ylabel('Accuracy (%)', fontsize=18)
axs[1].legend()
axs[1].grid(axis='y', linestyle='--')

fig.suptitle("Random Forest Model - Accuracy vs Weights", fontsize=18)
fig.savefig('fig/rf_model.png', format = 'png')
plt.savefig('../docs/assets/images/rf_model.png', format = 'png')

```



The best tree_depth is 5.

```
In [39]: tree_depth = 5
n_trees = 50
RF_model = RandomForestClassifier(n_estimators=n_trees,
                                max_depth=tree_depth,
                                class_weight = {0:1,1:20}).fit(X_train,
                                                                y_train)

performance(RF_model,X_train,X_test)
```

```
Rows: True Lables (0,1),
Columnns: Predicted Lables (0,1)
human bot
train:
[[575 107]
 [  0 118]]
test:
[[134  35]
 [  2  29]]
Training Set total accuracy: 86.625%
Training Set non-bot accuracy: 84.3108504398827%
Training Set bot accuracy: 100.0%
Test Set accuracy: 81.5%
Test Set non-bot accuracy: 79.28994082840237%
Test Set bot accuracy: 93.54838709677419%
```

bagging decision tree


```

In [40]: def perfromance_bagging():
    train_score = accuracy_score(y_train, y_pred_train)*100
    train_matrix = confusion_matrix(y_train, y_pred_train)
    test_score = accuracy_score(y_test, y_pred_test)*100
    test_matrix = confusion_matrix(y_test, y_pred_test)

    train_ture_negative = train_matrix[0][0]\
                        /(train_matrix[0][0]+train_matrix[0][1])*100
    train_ture_positive = train_matrix[1][1]\
                        /(train_matrix[1][0]+train_matrix[1][1])*100

    test_ture_negative = test_matrix[0][0]\
                        /(test_matrix[0][0]+test_matrix[0][1])*100
    test_ture_positive = test_matrix[1][1]\
                        /(test_matrix[1][0]+test_matrix[1][1])*100

    print()
    print("Rows: True Lables (0,1), \nColumnns: Predicted Lables (0,1)")
    print("human", "bot")

    print("train:")
    print(train_matrix)
    print("test:")
    print(test_matrix)

    print("Training Set total accuracy:",str(train_score)+'%')
    print("Training Set non-bot accuracy:",str(train_ture_negative)+'%')

    print("Training Set bot accuracy:",str(train_ture_positive)+'%')

    print("Test Set accuracy:",str(test_score)+'%')
    print("Test Set non-bot accuracy:",str(test_ture_negative)+'%')
    print("Test Set bot accuracy:",str(test_ture_positive)+'%')

```

```

In [41]: #Select overfit tree depth
         tree_depth = 10
         #Set number of trees
         n_trees = 11

         #Conduct bootstrapping and fit models to the data
         np.random.seed(0)
         bagging_train = np.zeros((X_train.shape[0], n_trees))
         bagging_test = np.zeros((X_test.shape[0], n_trees))
         bagging_models = []
         for i in range(n_trees):
             bootstrapped_X, bootstrapped_y = resample(X_train, y_train)
             fitted_model = DecisionTreeClassifier(max_depth=tree_depth,
                                                    class_weight = {0:1,1:1})\
                                                    .fit(bootstrapped_X,
                                                            bootstrapped_y)

             bagging_models.append(fitted_model)
             bagging_train[:,i] = fitted_model.predict(X_train)
             bagging_test[:,i] = fitted_model.predict(X_test)

         #Get Predictions across all models
         y_pred_train = np.mean(bagging_train, axis=1) > .5
         y_pred_test = np.mean(bagging_test, axis=1) > .5

```

```

In [42]: perfromance_bagging()

Rows: True Lables (0,1),
Columnns: Predicted Lables (0,1)
human bot
train:
[[682   0]
 [ 21  97]]
test:
[[158  11]
 [ 12  19]]
Training Set total accuracy: 97.375%
Training Set non-bot accuracy: 100.0%
Training Set bot accuracy: 82.20338983050848%
Test Set accuracy: 88.5%
Test Set non-bot accuracy: 93.49112426035504%
Test Set bot accuracy: 61.29032258064516%

```



```

In [45]: fig, axs = plt.subplots(1, 2, figsize=(15, 5))

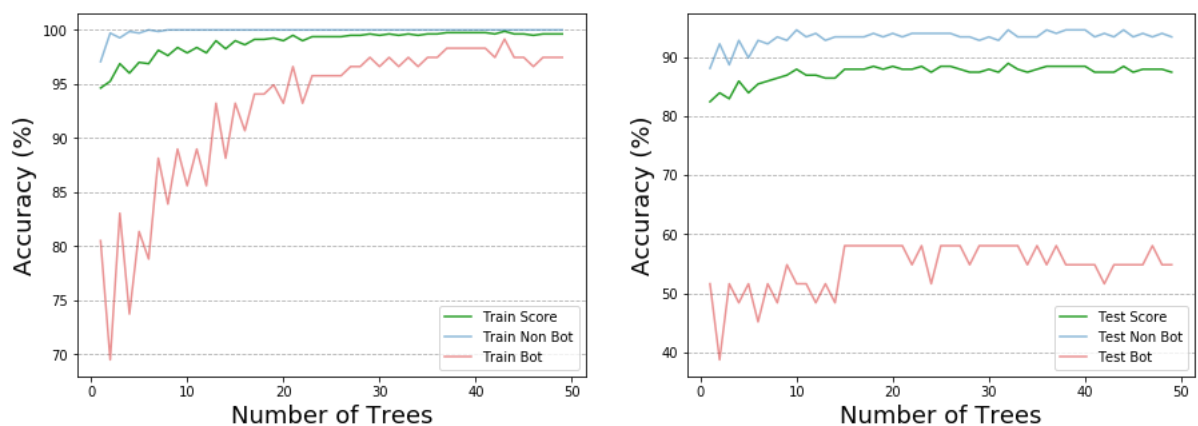
axs[0].plot(bagging_n_tree_accuracy.iloc[:,0],
            bagging_n_tree_accuracy.iloc[:,1],
            label='Train Score', color='C2')
axs[0].plot(bagging_n_tree_accuracy.iloc[:,0],
            bagging_n_tree_accuracy.iloc[:,2],
            label='Train Non Bot', color='C0', alpha=0.5)
axs[0].plot(bagging_n_tree_accuracy.iloc[:,0],
            bagging_n_tree_accuracy.iloc[:,3],
            label='Train Bot', color='C3', alpha=0.5)
axs[0].set_xlabel('Number of Trees', fontsize=18)
axs[0].set_ylabel('Accuracy (%)', fontsize=18)
axs[0].legend()
axs[0].grid(axis='y', linestyle='--')

axs[1].plot(bagging_n_tree_accuracy.iloc[:,0],
            bagging_n_tree_accuracy.iloc[:,4],
            label='Test Score', color='C2')
axs[1].plot(bagging_n_tree_accuracy.iloc[:,0],
            bagging_n_tree_accuracy.iloc[:,5],
            label='Test Non Bot', color='C0', alpha=0.5)
axs[1].plot(bagging_n_tree_accuracy.iloc[:,0],
            bagging_n_tree_accuracy.iloc[:,6],
            label='Test Bot', color='C3', alpha=0.5)
axs[1].set_xlabel('Number of Trees', fontsize=18)
axs[1].set_ylabel('Accuracy (%)', fontsize=18)
axs[1].legend()
axs[1].grid(axis='y', linestyle='--')

fig.suptitle("Bagging Model - Accuracy (%) vs Number of Trees",
            fontsize=16)
fig.savefig('../docs/assets/images/bagging_model.png', format = 'png')
fig.savefig('fig/bagging_model.png', format = 'png')

```

Bagging Model - Accuracy (%) vs Number of Trees



```
In [46]: def bagging_tuning_cut_off(cut_off):
#Conduct bootstrapping and fit models to the date
    np.random.seed(0)
    bagging_train = np.zeros((X_train.shape[0], n_trees))
    bagging_test = np.zeros((X_test.shape[0], n_trees))
    bagging_models = []
    for i in range(n_trees):
        bootstrapped_X, bootstrapped_y = resample(X_train, y_train)
        fitted_model = DecisionTreeClassifier(max_depth=tree_depth,
                                                class_weight = {0:1,1:2})\
                                                .fit(bootstrapped_X,
                                                        bootstrapped_y)

        bagging_models.append(fitted_model)
        bagging_train[:,i] = fitted_model.predict(X_train)
        bagging_test[:,i] = fitted_model.predict(X_test)

#Get Predictions across all models
    y_pred_train = np.mean(bagging_train, axis=1) > cut_off
    y_pred_test = np.mean(bagging_test, axis=1) > cut_off

    train_score = accuracy_score(y_train, y_pred_train)*100
    train_matrix = confusion_matrix(y_train, y_pred_train)
    test_score = accuracy_score(y_test, y_pred_test)*100
    test_matrix = confusion_matrix(y_test, y_pred_test)

    train_ture_negative = train_matrix[0][0]\
                        /(train_matrix[0][0]+train_matrix[0][1])*100
    train_ture_positive = train_matrix[1][1]\
                        /(train_matrix[1][0]+train_matrix[1][1])*100

    test_ture_negative = test_matrix[0][0]\
                        /(test_matrix[0][0]+test_matrix[0][1])*100
    test_ture_positive = test_matrix[1][1]\
                        /(test_matrix[1][0]+test_matrix[1][1])*100

    accuracy = [cut_off, train_score, train_ture_negative,
                train_ture_positive, test_score, test_ture_negative,
                test_ture_positive]
    return accuracy
```

[illegible]

```

In [48]: fig, axs = plt.subplots(1, 2, figsize=(15, 5))


axs[0].plot(bagging_cut_off_accuracy.iloc[:,0],
            bagging_cut_off_accuracy.iloc[:,1],
            label='Train Score', color='C2')
axs[0].plot(bagging_cut_off_accuracy.iloc[:,0],
            bagging_cut_off_accuracy.iloc[:,2],
            label='Train Non Bot', color='C0', alpha=0.5)
axs[0].plot(bagging_cut_off_accuracy.iloc[:,0],
            bagging_cut_off_accuracy.iloc[:,3],
            label='Train Bot', color='C3', alpha=0.5)
axs[0].set_xlabel('Cut Off', fontsize=18)
axs[0].set_ylabel('Accuracy (%)', fontsize=18)
axs[0].legend()
axs[0].grid(axis='y', linestyle='--')

axs[1].plot(bagging_cut_off_accuracy.iloc[:,0],
            bagging_cut_off_accuracy.iloc[:,4],
            label='Test Score', color='C2')
axs[1].plot(bagging_cut_off_accuracy.iloc[:,0],
            bagging_cut_off_accuracy.iloc[:,5],
            label='Test Non Bot', color='C0', alpha=0.5)
axs[1].plot(bagging_cut_off_accuracy.iloc[:,0],
            bagging_cut_off_accuracy.iloc[:,6],
            label='Test Bot', color='C3', alpha=0.5)
axs[1].set_xlabel('Cut Off', fontsize=18)
axs[1].set_ylabel('Accuracy (%)', fontsize=18)
axs[1].legend()
axs[1].grid(axis='y', linestyle='--')

fig.suptitle("Bagging Model - Accuracy vs Cut Off", fontsize=18)
fig.savefig('../docs/assets/images/bagging_model_cut_off.png',
            format = 'png')
fig.savefig('fig/bagging_model_cut_off.png', format = 'png')

```

Bagging Model - Accuracy vs Cut Off



```

In [49]: #Select overfit tree depth
tree_depth = 10
#Set number of trees
n_trees = 51

#Conduct bootstrapping and fit models to the data
np.random.seed(0)
bagging_train = np.zeros((X_train.shape[0], n_trees))
bagging_test = np.zeros((X_test.shape[0], n_trees))
bagging_models = []
for i in range(n_trees):
    bootstrapped_X, bootstrapped_y = resample(X_train, y_train)
    fitted_model = DecisionTreeClassifier(max_depth=tree_depth,
                                         class_weight = {0:1,1:1})\
                                         .fit(bootstrapped_X,
                                             bootstrapped_y)

    bagging_models.append(fitted_model)
    bagging_train[:,i] = fitted_model.predict(X_train)
    bagging_test[:,i] = fitted_model.predict(X_test)

#Get Predictions across all models
y_pred_train = np.mean(bagging_train, axis=1) > .2
y_pred_test = np.mean(bagging_test, axis=1) > .2

```

```

In [50]: perfromance_bagging()

```

```

Rows: True Lables (0,1),
Columnns: Predicted Lables (0,1)
human bot
train:
[[663  19]
 [  1 117]]
test:
[[145  24]
 [  5  26]]
Training Set total accuracy: 97.5%
Training Set non-bot accuracy: 97.21407624633432%
Training Set bot accuracy: 99.15254237288136%
Test Set accuracy: 85.5%
Test Set non-bot accuracy: 85.79881656804734%
Test Set bot accuracy: 83.87096774193549%

```

website frequency

```
In [51]: predictions=pd.DataFrame(columns=['user_id',
                                         'user_name',
                                         'scores_universal',
                                         'is_bot_by_botometer',
                                         'mean_links_to_twitter',
                                         'mean_links_to_top_social_media',
                                         'mean_links_to_top_digital_media',
                                         'mean_links_to_top_news',
                                         'mean_links_to_top_products_services',
                                         'mean_links_to_top_celebrities',
                                         'mean_links_to_top_organizations',
                                         'mean_links_to_top_sports',
                                         'mean_links_to_top_adult',
                                         'prediction_lg',
                                         'prediction_RF',
                                         'prediction_bagging_tree'])

is_bot_by_botometer=y_data
```

```
In [52]: predictions['user_id'] = users_summary_df['user.id_str']
predictions['user_name'] = users_summary_df['name']
predictions['scores_universal'] = users_summary_df['scores.universal']
predictions['is_bot_by_botometer'] = is_bot_by_botometer

for i in ['mean_links_to_twitter',
          'mean_links_to_top_social_media',
          'mean_links_to_top_digital_media',
          'mean_links_to_top_news',
          'mean_links_to_top_products_services',
          'mean_links_to_top_celebrities',
          'mean_links_to_top_organizations',
          'mean_links_to_top_sports',
          'mean_links_to_top_adult']:
    predictions[i] = users_summary_df[i]
predictions.index = range(1000)
```

```
In [53]: train_index = X_train.index.tolist()
test_index = X_test.index.tolist()
index = train_index + test_index
y_pred_bag = pd.DataFrame(y_pred_train.tolist() + y_pred_test.tolist())
y_pred_bag.columns = ['pred_bag']
y_pred_bag.index = index
y_pred_bag = y_pred_bag.sort_index()
y_pred_bag = pd.get_dummies(y_pred_bag['pred_bag'], drop_first=True)
y_pred_bag.columns = ['pred_bag']
```

```
In [54]: predictions['prediction_lg'] = LR_model.predict(X_data)
predictions['prediction_RF'] = RF_model.predict(X_data)
predictions['prediction_bagging_tree'] = y_pred_bag['pred_bag']
```



```
In [55]: vote = []
for i in range(len(predictions)):
    if np.mean(predictions.iloc[i,-4:-1]) > 0.5:
        vote = vote + [1]
    else: vote = vote + [0]
```

```
In [56]: vote = []
for i in range(len(predictions)):
    if np.mean(predictions.iloc[i,-4:-1]) > 0.5:
        vote = vote + [1]
    else: vote = vote + [0]
```

```
In [57]: predict_bot_index=predictions.index[predictions['prediction_RF']==1]
predict_non_bot_index=predictions.index[predictions['prediction_RF']==0]
```

```
In [58]: total_website_sum = []
bot_website_sum = []
ratio = []

for i in ['mean_links_to_twitter',
          'mean_links_to_top_social_media',
          'mean_links_to_top_digital_media',
          'mean_links_to_top_news',
          'mean_links_to_top_products_services',
          'mean_links_to_top_celebrities',
          'mean_links_to_top_organizations',
          'mean_links_to_top_sports',
          'mean_links_to_top_adult']:
    total_website_sum = total_website_sum + [sum(predictions[i])]
    bot_website_sum = bot_website_sum \
        + [sum(predictions[i][predict_bot_index])]
    ratio = ratio + [sum(predictions[i][predict_bot_index])\
                    /sum(predictions[i])]

total_website_sum = total_website_sum + [sum(total_website_sum)]
bot_website_sum = bot_website_sum + [sum(bot_website_sum)]
ratio = ratio + [sum(bot_website_sum)/sum(total_website_sum)]
```

```
In [59]: bot_website_ratio = pd.DataFrame(ratio)
bot_website_ratio.index = ['mean_links_to_twitter',
                           'mean_links_to_top_social_media',
                           'mean_links_to_top_digital_media',
                           'mean_links_to_top_news',
                           'mean_links_to_top_products_services',
                           'mean_links_to_top_celebrities',
                           'mean_links_to_top_organizations',
                           'mean_links_to_top_sports',
                           'mean_links_to_top_adult',
                           'Total']
```

```
In [60]: bot_website_ratio.iloc[0:9].sort_values(by=[0])
```

Out[60]:

| | 0 |
|-------------------------------------|----------|
| mean_links_to_top_sports | 0.090909 |
| mean_links_to_twitter | 0.208805 |
| mean_links_to_top_celebrities | 0.393511 |
| mean_links_to_top_digital_media | 0.454819 |
| mean_links_to_top_social_media | 0.555639 |
| mean_links_to_top_news | 0.593913 |
| mean_links_to_top_products_services | 0.621033 |
| mean_links_to_top_organizations | 0.967320 |
| mean_links_to_top_adult | 1.000000 |

```
In [61]: bot_website_ratio_each = bot_website_ratio.iloc[0:9].sort_values(by=[0])
bot_website_ratio_total= bot_website_ratio.iloc[9:10]
bot_website_ratio_total[0]
```

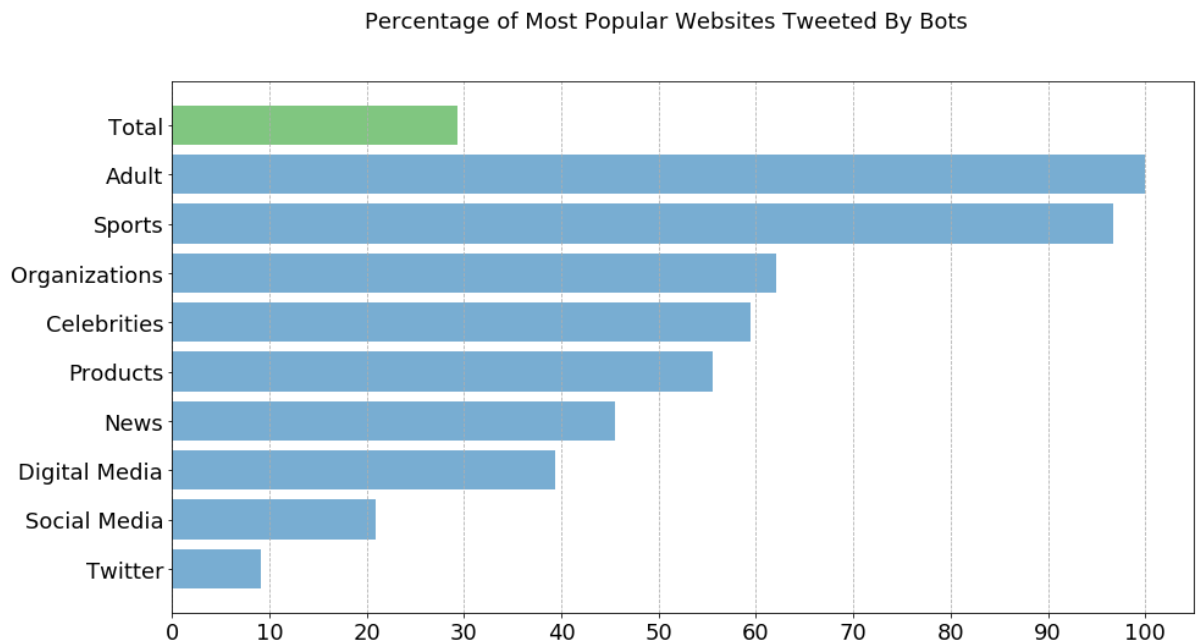
Out[61]: Total 0.293066
Name: 0, dtype: float64

```

In [62]: plt.figure(figsize=(15, 8))
plt.barh(range(9),bot_website_ratio_each[0] * 100, alpha = 0.6,
         color = 'C0')
plt.barh(9,bot_website_ratio_total[0] * 100, alpha = 0.6,
         color = 'C2')
plt.yticks(range(10), ['Twitter',
                        'Social Media', 'Digital Media',
                        'News', 'Products',
                        'Celebrities', 'Organizations',
                        'Sports', 'Adult','Total'], fontsize = 18)
plt.xticks(range(0,110,10),fontsize=18)
plt.grid(axis='x', linestyle='--')

plt.suptitle("Percentage of Most Popular Websites Tweeted By Bots",
            fontsize=18)
plt.savefig('fig/websites.png', format = 'png')
plt.savefig('../docs/assets/images/websites.png',
            format = 'png')
plt.show()

```



In []: